

Speeding up evaluation of structures for the Angry Birds Game

First Author Name¹^a, Second Author Name¹^b and Third Author Name²^c

¹Place Holder, XYZ University, My Street, MyTown, MyCountry

²Place Holder, Main University, MySecondTown, MyCountry
{f_author, s_author}@ips.xyz.edu, t_author@dc.mu.edu

Keywords: Search-Based Procedural Content Generator, Evolutionary algorithm, Game development, Angry Birds, Level generation

Abstract: In this work, we present an original method based on evolutionary algorithms for generating basic structures for the physics-based game Angry Birds, with the ultimate objective of creating Angry Birds levels with the minimum number of constraints. We set out to evolve free-form structures, and this means searching in a larger space. In this paper, we test how using a physics engine enables us to evaluate much more levels than a game engine simulation. In order to do this, we compare the results of experiments using both types of simulators and propose fitness functions accordingly. Results show the execution time drastically drops from 5 hours to less than 20 minutes on average.

1 INTRODUCTION


Angry Birds is a multiplatform video game created in 2009 by the Rovio Entertainment Corporation. Each level of the game consists of a collection of structures made out of blocks in which comic *pig* characters are hiding; the player has to fire from a slingshot different *bird* characters, each having different abilities and moods. The objective of the game is to destroy all the pigs by knocking down the structures or just hitting them directly. The game relies on gravity to create interesting puzzles by closely resembling the dynamics of real-world structures. It has been approached in different ways from the computational intelligence community; in this paper we are interested in generating levels that are playable and interesting. This has been proposed as a competition in several game-related conferences (Stephenson et al., 2018; Khalifa, 2018), but could be also interesting from the perspective of generating levels to train or evaluate game-playing bots, as was done by Zafar et al. in (Zafar et al., 2019).


Search-based Procedural Content Generation (SBPCG) is a type of a *generate-and-test* approach to PCG which is usually tackled with Evolutionary Algorithms(EA) (Togelius et al., 2010). The chal-


lenges faced by SBPCG are not far from those found in EAs; since they are search methods, they can be a good option to implement this kind of systems. In our case, as a first step to generate Angry Birds levels, we will generate free-form self-standing structures using the Angry Birds basic blocks. The main intention with this approach is to eventually generate structures that can host *pigs* and that achieve aesthetic quality mainly through variability and the fact that they are not cookie-cutter repetitions of the same basic structure with some small variations; at the same time, it becomes an interesting and challenging problem from the optimization point of view: using some basic building blocks and a simulated gravity, to be able to generate structures that do not collapse. That is the main objective of this line of work.

The primary objective of this paper, however, is to use a fitness function that is able to evaluate structures without needing the simulator, advancing the state of the art with respect to our previous paper (Calle et al., 2019), where we needed to use the simulator itself to check gravity, which incurred in much overhead, up to several seconds.

In this paper, we will use what (Togelius et al., 2010) calls *direct fitness function*, this function computes a score from measurable features of the generated content. However, this fitness function is a time-consuming task since it involves the generation of a graphic representation of the structure and the simulation of the falling motion. If we have to evaluate

^a <https://orcid.org/0000-0000-0000-0000>

^b <https://orcid.org/0000-0000-0000-0000>

^c <https://orcid.org/0000-0000-0000-0000>

every single individual in the population, we will not be able to cover enough of the free-form search space to find a good enough solution. So we must minimize the actual number of structures that are simulated by applying heuristics to the data structure and assigning it a fitness even before the simulation.

This paper is organized as follows: in the next section, we present the state of the art in this type of level generation, as well as its relation to the problem of generating structurally sound structures. Our proposed method for generating Angry Birds levels is described next in Section 3. Next experiments are presented in section 4. We present our conclusions in 5.

2 Background and State of the Art

PCG used for video game levels is relevant in international artificial intelligence competitions, such as Super Mario Level Generation (Shaker et al., 2012), General Video Games (Khalifa, 2018; Khalifa et al., 2016), or recently, for Angry Birds (Stephenson et al., 2018). This work is related to works presented in previous editions of the competition. The focus of the latest edition (Stephenson et al., 2018) was on finding entertaining levels. Fun was the main factor in the evaluation of proposals; secondary factors were creativity and difficulty. Six entries participated, of which J. Yuxuan et al., were able to generate random quotations with different components of a level; J. Xu et al. generated levels that look like pixel images. A third approach (by C. Kocaogullar) translated music patterns to generate structures. The winner was *Iratus Aves*, a new iteration of the work by M. Stephenson and J. Renz (Stephenson and Renz, 2017; Stephenson and Renz, 2016), which follows a *constructive method*. In this work, the likelihood of selecting a certain block is given by a probability table, which was tuned using an optimization method. Blocks are then stacked following a tree structure.

A constructive method ensures local stability, but not global stability, which must be evaluated once the whole structure is completed. One problem with this and other constructive approaches is that the variety of structures created is going to be relatively small; monotony leads to boredom, decreasing playability. On the other hand, generated structures are guaranteed to be structurally sound, and constructive approaches are generally faster than search-based procedures.

An alternative to deal with these limitations is to follow a *Search-based approach*. Thus, Lucas Ferreira and Claudio Toledo (Ferreira and Toledo, 2014)

presented a solution using a Genetic Algorithm (GA) and a clone of Angry Birds named *Science Birds* developed to evaluate the levels. In this GA, levels correspond to individuals in a population, each individual has a chromosome represented by an array of lists. Each list is a sequence of blocks, pigs and predefined compound blocks, using an identification number. Each list is then placed as a stack of elements in the game. A level has several such stacks. This representation also includes the distances between stacks. The population is initialized randomly following a probability table defining the likelihood of a certain element being placed in a certain position within a stack. This implies that a column or stack shape is chosen beforehand, once again ensuring stability, but decreasing playability by generating structures whose only differences are which blocks are placed on top of which.

Levels are evaluated executing them in the simulator and checking their average stability, considering the speed of every block when erected – which must tend to be zero when having a stable structure –. The authors designed specialized crossover and mutation operators, aiming to maintain the consistency of newly generated solutions.

The approach proposed by Stephenson et al. (Stephenson and Renz, 2019) is based on agents, and is focused on offering custom experience for specific players.

However, this work proposes a different approach: *free-form evolution*. If we look outside the domain of game development and focus on structural optimization in architecture, there are several proposals using search-based algorithms. We can find a metaheuristic called Cuckoo Search (Gandomi et al., 2013) which performance was tested with structural optimization. However, this optimization is heavily parametrized and we are looking for the evolution of structures that do not follow a predefined pattern. Another approach for structure design is using Generative Grammatical Encodings (Hornby and Pollack, 2001) where L-system and its production rules are considered individuals. This method increases the number of generated patterns, but still restrains the formation of disjoint structures, for instance, a defensive tower before a simple pigsty in our context.

We aim at following a realistic structure generation approach, without constraining it to a fixed form, thus advancing the state of the art by allowing the creation of Angry Birds levels having any arrangement. The next section will describe how we characterize this problem and our approach to it. In our previous paper (Calle et al., 2019) we explored this approach and found as one of its shortcomings the fact that the

evaluation of generated structures through the simulator was lengthy and didn't leave the algorithm enough time to explore the search space. In this paper, we try to tackle that problem, as well as take additional steps to increase the complexity of generated structures.

3 Problem Description

In a previous work, we used *Science Birds* (Ferreira and Toledo, 2014) as a starting point. Developed by Lucas Ferreira and Claudio Toledo, Science Birds has become the main open source Angry Birds simulator. However, we needed to modify the code in order to produce an output with the additional data needed to automate that work. The code is available in the GitHub repository <https://github.com/Laucalle/ScienceBirds>. The additional data contains the position and the average magnitude of the velocity of each block that remains after the simulation. One advantage is that the whole experiment can run without user intervention. Another advantage is that the amount of time spent on the simulation of each level is minimized to some extent. Reducing the simulation time not only increases the number of evaluations by a certain range of time, but it is also a factor in competitions, where there is a fixed time limit for participants.

To further decrease simulation time, in this paper, we use the Box2D (<https://box2d.org>) physics engine. This engine, written in C++, was initially used for the actual Angry Birds game. By just simulating the positioning of objects in memory, we can avoid launching the whole game, using bitmaps and screen rendering, which adds overhead to the process of fitness evaluation. Even if in this Box2D simulation we do not have the resistance of the blocks implemented, we can test the stability of level much more efficiently, since there is no overhead in computing things unrelated to Physics, such as the GUI. If we do not have block resistance in a simulation, blocks will not be destroyed when they hit each other; in that case, the fitness function should not take into account before and remaining blocks.

Once we chose the simulator that is going to be used to evaluate the fitness of a level, we must design the fitness function. As obvious as it might seem, the main feature of a sound level is that it is not in motion, so it seems reasonable to evaluate its complete stillness as opposed to its speed. We must consider every single block in doing this.

$$fitness_{ind} = \frac{1}{|V|} \sum_{i=0}^b V_i + P_{broken} \cdot (b - |V|)$$

When using Science Birds, the average magnitude of velocity is provided for each block. We note this as V , with $|V|$ being the cardinality of the set. The number of blocks in an individual is b , and it can differ from the cardinality of V since we do not track collapsed blocks. The number of broken blocks is $b - |V|$, and it is multiplied by a penalization factor since a level whose blocks break without user interaction would not be considered valid. Blocks can be broken when they fall from a certain height or collide with another object. We set the penalization factor P_{broken} to 100 since objects in a level do not usually reach that velocity. The goal of this evaluation is then to separate non-valid levels from potentially good ones.

In the experiment, presented we changed this fitness function to the function shown below, after observing the results for the previous experiments 4:

$$fitness_{indV2} = \max(V)$$

As we said before, simulating a level is a highly time-consuming task, much more when we simulate the whole game, it is in the order of seconds, which makes it almost unfeasible for our purpose. Next, we propose a method for evaluation, in which not all levels need are simulated.

Some situations can indicate that a level has a very slim probability of being valid. For instance, a block begins suspended in a position to far from the ground, or there are blocks with an overlapping position. If this is the case, then we skipped the simulation of the level.

When having a structure where the object closest to the ground is far above, it will likely collapse from the impact along with all the other blocks above. For this reason, we will not be simulating levels that have all their blocks higher than a certain threshold. The threshold used is 0.1 in-game units, and the penalty applied to the distance is 10:

$$f_{distance} = \begin{cases} P_{distance} \cdot D_{lowest}, & \text{if } D_{lowest} > threshold \\ 0, & \text{otherwise} \end{cases}$$

The other factor is the number of overlapping blocks. To determine if two convex shapes intersect, we can use the separating axis theorem (Ericson, 2004) used in game development for collision detection. A level with blocks that occupy the same space is not likely to be stable, as the Unity Engine underlying the simulator will solve the issue moving the blocks until there is no collision. Unity implements this behavior, and as it is proprietary software, it is not possible to know or change what it does. So, a penalty is also applied and the level is not simulated either.

For this case it is $f_{\text{overlapping}} = P_{\text{overlapping}} \cdot N_{\text{overlapping}}$ where the first factor is a penalty set to 10 and the second is the number of blocks overlapping with each other.

In some of the late experiments, we will substitute the f_{distance} with the gap in the Y-axis. We then project all blocks on the Y-axis and calculate the range of values in the Y coordinate that are not inside the feasible range. This gap is treated the same way as f_{distance} (same penalization and threshold) so we call it $f_{Y\text{-axis}}$:

$$f_{Y\text{-axis}} = \begin{cases} P_{\text{distance}} \cdot P_{Y\text{-axis}}, & \text{if } P_{Y\text{-axis}} > \text{threshold} \\ 0, & \text{otherwise} \end{cases}$$

If both f_{distance} and $f_{\text{overlapping}}$ are 0 then the level is suitable for simulation and fitness is calculated as $\text{fitness}_{\text{ind}}$. We could consider this approach as an *overpenalization* but exploring unfeasible regions entails a serious overhead that we need to minimize (Runarsson and Yao, 2003). On the other hand, levels with multiple or even all blocks broken during the simulation are not feasible either but in this case, running the simulation is necessary. In this last case, penalization does not prevent the algorithm from exploring that region.

Since one of the perceived benefits of SBPCG is the expressiveness and variability, it seems reasonable to use a flexible representation. We will design the GA to allow a less directed search than previous proposals while keeping the representation simple.

Individuals are composed of a list of blocks; we do not consider, TNT boxes, or pigs in this paper since we are focused only on the generation of structures. These building blocks have the following attributes:

- Type: there are only eight basic blocks that can be placed in the level with different shapes and sizes; they are represented by an integer between 0 and 7.
- Position: x and y coordinates from the centre of the block. Values are in game units and are represented as floating point numbers.
- Rotation: rotation of the basic block in degrees. Only four different rotation angles are considered: 0, 45, 90 or 135 degrees represented as integers between 0 and 3.
- Material: there are three types *wood*, *metal* and *glass*, which determine the durability of the block. However, this does not affect their stability, so we only use *wood* material for now.

Using this representation a gene representing a single block will be formed by two integers (type and

position) and two floating point numbers (x and y coordinates).

Individuals are a collection of genes, in the same way a level is a collection of building blocks. The number of blocks is variable and the order in which they are listed is not important.

We store the fitness of the worst individual that has been tested in-game so that the value of not tested levels is always above—it is a minimization problem—the in-game tested levels; the starting point for fitness of such individuals is the worst in-game score.

This penalization is calculated using the distance of the lowest block to the ground, which can be easily obtained, and the number of blocks that overlap. This requires a bit more of computation, so it will be stored and set in the initialization of the individual. When a gene is modified, the number of overlapping blocks is recalculated for that specific change.

Considering all of the above, the chromosome object is composed by:

- A non-fixed list of genes.
- A fitness value.
- A penalty (set to `False` for in-game evaluated levels).
- The number of overlapping blocks (calculated).

Initialization is done randomly, with each individual having a random number of genes, which are initialized by several methods:

- Random: selects a random number for each attribute of the gene.
- Non-overlapping: also selects a random number but the gene is only added to the chromosome if it does not overlap with an already existing gene.
- Discrete: selects a random number for type and rotation, but the position must be multiple of the dimensions of the smallest block (blocks will be aligned).
- Discrete non-overlapping: it combines the second and third initialization methods.
- Discrete with a set of pre-configured blocks: first it includes a set of blocks, and then adds blocks following the third method until it reaches the desired number of blocks. The configurations used are the compound blocks found in (Ferreira and Toledo, 2014).

Candidates for reproduction are selected using tournaments. Two individuals are chosen from the population and the best will be a parent in this generation. This is repeated until a certain percentage of pairs have been reached. It is important to note that

individuals chosen are not removed from the population and therefore they can appear several times in the list of parents.

Once the parents have been selected, we chose from two different methods of combination:

- **Sample Crossover:** gives a single individual per parent pair. It takes all genes from both parents—excluding genes that are repeated—and randomly takes a number of them to create the new individual. The number of blocks is the minimum between the maximum number of blocks allowed, the mean of the two parent individuals and the number of distinct genes.
- **Common Blocks:** produces two individuals. The common genes to both parents are passed on to both children. The remaining genes are randomly distributed to each child, half to one and half to the other.

There are four different types of mutation:

- **Rotation:** rotation is represented as an integer (it is discretized), so it adds or subtracts one to the current value.
- **Type:** similarly to rotation mutation.
- **Position X:** a real value between 0 and 1—excluding 0—is added or subtracted from the value of the position X.
- **Position Y:** same as position X mutation, for position Y.

They are all applied to random members of the population.

The new generation is produced following an elitist strategy. Best individuals in both the old population and their offspring pass on to the next generation, maintaining the size of the population.

The information that describes a level can be too complex to have a binary representation as pure genetic algorithms suggest, so the framework used should be flexible enough to support complex data structures. This prevented us from using other frameworks and therefore a new framework was implemented. The source code is open source and can be found again in GitHub at <https://github.com/Laucalle/AngryBirdsLevelGenerator>.

In order to assess the proposed methods and verify if they meet our objective, we performed a series of experiments presented in the next.

4 Experiments and Results

We set out to evolve free-form structures, and this means searching in a larger space. In this experiment,

	Time(h)	σ	G	σ
E1	0.89	(0.59)	100.0	(0)
E2	1.002	(1.97)	155.087	(240.56)
E3	1.76	(0.6)	76.625	(42.3)
E4	5.03	(1.46)	365.929	(158.09)
E5	0.099	(0.1)	121.2	(96.89)
E6	0.788	(0.124)	1000.0	(0)

Table 1: Summary of the execution of the last generation in 15-20 runs for each experiment. 40 runs for E5 and E6. G: number of generations, E: experiment number

we test how using a Physics engine enables us to evaluate much more levels than a game engine simulation. In order to do this, we compare the results of experiments with both simulators. Tables 1 and 2 show an overview of the results. Experiments **E1** to **E4** were implemented using a game engine simulator while **E5** and **E6** a physics engine.

4.1 Removing game and penalizing gaps in Y-axis

The main problem with the previous experiments (**E1-E4**) was the time needed to load the Science Birds simulation environment to run the levels, which needed several seconds for loading and obtaining results. So the main objective of this experiment was to find a way to get rid of the in-game simulation. In order to do that, we will use the Box2D (Catto, 2011) Physics engine we mentioned earlier.

Since game physics do not usually resemble real world physics we adjusted the parameters so this simulation and the game behave in the same way. As we can see in table 1 the execution time drastically drops from 5 hours to less than 20 minutes on average, even when running more generations in the process.

This opened the way for performing more operations for the individuals. In this case we chose to penalize not only the distance to the ground but also the *gaps* in the Y-axis, which will make objects drop and maybe break. This will encourage individuals to grow vertically and not only horizontally like in previous experiments. This changes the fitness function, so we will have to compare by the actual obtained structure, one of which is shown in Figure 1.

In general, this penalization of gaps creates a faster path to stable structures. Still, this path leads to mostly flat structures with some higher block in unstable positions, which are structurally solid, but not interesting. One of the stable results is shown in figure 1.

	Best	σ	Avg	σ	Worst	σ
E1	61.334	(133.02)	383.701	(106.14)	510.515	(133.04)
E2	110.66	(142.21)	327.547	(238.33)	367.895	(260.83)
E3	0.0015	(0.003)	0.54	(0.24)	0.828	(0.34)
E4	0.0018	(0.003)	0.203	(0.068)	0.2997	(0.1)
E5	1.249	(1.257)	1.276	(1.231)	1.288	(1.219)
E6	1.031	(0.853)	1.27	(0.834)	1.328	(0.819)

Table 2: Summary of the results of the last generation in 15-20 runs for each experiment. 40 runs for E5 and E6.




Figure 1: One of the solutions from 4.1



Figure 2: One of the solutions from 4.2

4.2 Changing the evaluation function

Observing results from the previous experiment we realized that what evolution found was that laying many blocks on the ground was enough to get a high fitness: the average speed was decreased and it will place unstable blocks to cover gaps in Y-axis. In order to correct this behaviour we changed the fitness function to take into account the fastest moving object:

$$fitness_{indv2} = \max(V)$$

Additionally, we initialized levels including one of a list of pre configured blocks in addition to the random initialization used until now.

This makes the fitness value depends on just one gene, although it can be a different one each time. The improvement of solutions to find acceptable ones slowed down again, with a different fitness function we cannot compare the fitness value with the rest of the experiments.

5 Conclusions and Future Work

This paper was developed with the main objective of implementing a competitive entry for the Angry Birds level generation contest, and a set of sub-objectives in mind: exploring the expressiveness and variability of SBPCG with evolutionary techniques, adapting the game to extract data from execution and producing stable structures under gravity, in a simulated environment.

For this aim we have implemented an Evolutionary Algorithm able to optimize level structures to meet some criteria, being the main one the stability of the constructions. Considering this, the method studied was sufficiently general and flexible to draw some conclusions about the topic. SBPCG methods are a potential good solution to offline content generation but it requires a great amount of problem-specific knowledge. Like any other form of creative work, the biggest issue may be how to measure how good, creative or enjoyable is the piece. The more rules the author adds, expressiveness starts to get lost as the results are variations of the same idea. However, it is crystal clear from the experiments run in this paper that a lack of knowledge will lead to unexpected outcomes. In our case, the fitness function used the stability of the structure and only considered other features— overlapping blocks and distance to the ground— to ensure the levels would be valid. In experiment 4.1 and 4.2 height was also taken into account.

The second objective, adapting the game to extract data from execution, was certainly achieved. It was also a basic requirement to proceed with the rest of them. The game does provide the data, as long as the input is correctly structured. In order to conduct our experiments the Angry Birds simulator (Science Birds) has been adapted to our necessities, yielding now some other information required to evaluate the individuals of the implemented EA. As this was a bottleneck, experiments 4.1 and 4.2 used only a physics simulation, leaving behind the actual game. It would be interesting to obtain other kind of data from both simulations, such as the height of the structure and,

eventually, its resistance to bombardment by angry birds. However, this is left as future work.

Producing stable structures under gravity was the third objective and the closest to the ultimate one. That objective was effectively achieved, but the consequence of evolving in a path of minimum movement or maximum stability results in structures that are close to the ground and do not have more than a few floors, which could not be very exciting for the players. The main issue is, then, how we evaluate the levels. In fact this is a matter of how we define what an Angry Birds' level *is* and if that definition matches the fitness function. How this definition of level plays with the paths of evolution is also a problem. Since we define as level as a structure that is stable, evolution will maximize stability finding, as in the beginnings of architectural practice, squat structures that are neither aesthetically pleasing nor playable, although undeniably sturdy and stable. We have been successful in, evolving free form structures, to find these type. But once we get there, we need to go into a different evolution mode that takes into account several features as a multiobjective optimization problem: stability is the first, but we can sacrifice a bit of stability for height or some other aesthetic quality.

To conclude on an optimistic tone, this work provides an interesting insight into the SBPCG, through the completion—and failure—of the goals we set out to achieve at the beginning.

In order to improve the results of the method, different constraints could be expressed as multiple objectives. Overlapping blocks and velocity could be treated as minimization objectives and height as a maximization one.

If we pay attention at the stages of evolution in this work, there is also room for improvement in the genetic operators. For example, the initialization produces a small amount of valid individuals which suggested that an elitist strategy for selection would work best. However, new experiments will help to better balance exploration and exploitation. An interesting addition would be to add *building* operators that pile blocks on structures that are already stable.

6 Acknowledgements

This paper has been supported in part by DeepBio (TIN2017-85727-C4-2-P) from the Ministerio de Economía y Competitividad in Spain.

REFERENCES

- Calle, L., Merelo Guervós, J. J., García, A. M., and García Valdez, J. M. (2019). Free form evolution for angry birds level generation. In Kaufmann, P. and Castillo, P. A., editors, *Applications of Evolutionary Computation - 22nd International Conference, EvoApplications 2019, Held as Part of EvoStar 2019, Leipzig, Germany, April 24-26, 2019, Proceedings*, volume 11454 of *Lecture Notes in Computer Science*, pages 125–140. Springer.
- Catto, E. (2011). Box2D: A 2D Physics engine for games.
- Ericson, C. (2004). *Real-time collision detection*. CRC Press.
- Ferreira, L. and Toledo, C. (2014). A search-based approach for generating angry birds levels. In *Computational intelligence and games (cig), 2014 IEEE conference on*, pages 1–8. IEEE.
- Gandomi, A. H., Yang, X.-S., and Alavi, A. H. (2013). Cuckoo search algorithm: a metaheuristic approach to solve structural optimization problems. *Engineering with computers*, 29(1):17–35.
- Hornby, G. S. and Pollack, J. B. (2001). The advantages of generative grammatical encodings for physical design. In *Evolutionary Computation, 2001. Proceedings of the 2001 Congress on*, volume 1, pages 600–607. IEEE.
- Khalifa, A. (2018). The general videogame ai competition - level generation track. <https://github.com/GAIGResearch/gvgai/wiki/Tracks-Description/#level-generation-track>. Accessed: 2018-11-05.
- Khalifa, A., Perez-Liebana, D., Lucas, S. M., and Togelius, J. (2016). General video game level generation. In *Proceedings of the Genetic and Evolutionary Computation Conference 2016, GECCO '16*, pages 253–259, New York, NY, USA. ACM.
- Runarsson, T. P. and Yao, X. (2003). Evolutionary search and constraint violations. In *Evolutionary Computation, 2003. CEC'03. The 2003 Congress on*, volume 2, pages 1414–1419. IEEE.
- Shaker, N., Togelius, J., and Yannakakis, G. (2012). Mario ai championship - level generation track. <http://www.marioai.org/LevelGeneration>. Accessed: 2018-11-05.
- Stephenson, M. and Renz, J. (2016). Procedural generation of complex stable structures for angry birds levels. In *Computational Intelligence and Games (CIG), 2016 IEEE Conference on*, pages 1–8. IEEE.
- Stephenson, M. and Renz, J. (2017). Generating varied, stable and solvable levels for angry birds style physics games. In *Computational Intelligence and Games (CIG), 2017 IEEE Conference on*, pages 288–295. IEEE.
- Stephenson, M. and Renz, J. (2019). Agent-based adaptive level generation for dynamic difficulty adjustment in angry birds. *arXiv preprint arXiv:1902.02518*.
- Stephenson, M., Renz, J., Ferreira, L., and Togelius, J. (2018). 3rd angry birds level generation competition.

<https://project.dke.maastrichtuniversity.nl/cig2018/competitions/\#angrybirds>. Accessed: 2018-11-05.

- Togelius, J., Yannakakis, G. N., Stanley, K. O., and Browne, C. (2010). Search-based procedural content generation. In *European Conference on the Applications of Evolutionary Computation*, pages 141–150. Springer.
- Zafar, A., Hassan, S., et al. (2019). Corpus for angry birds level generation. In *2019 2nd International Conference on Computing, Mathematics and Engineering Technologies (iCoMET)*, pages 1–4. IEEE.